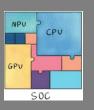


ECE4810J System-on-Chip Design



Topic 6

Reconfigurable SoC

Xinfei Guo xinfei.guo@sjtu.edu.cn

November 11th, 2024

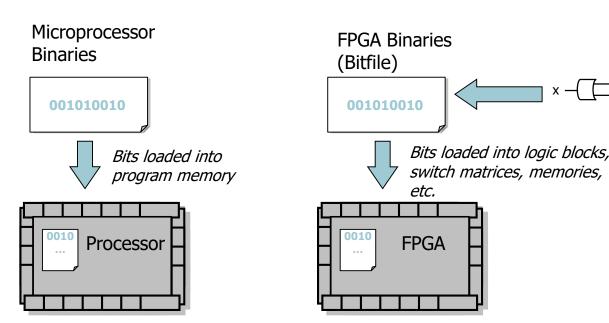


T7 learning goals

- Reconfigurable SoC
 - Why?
 - How?
 - What?

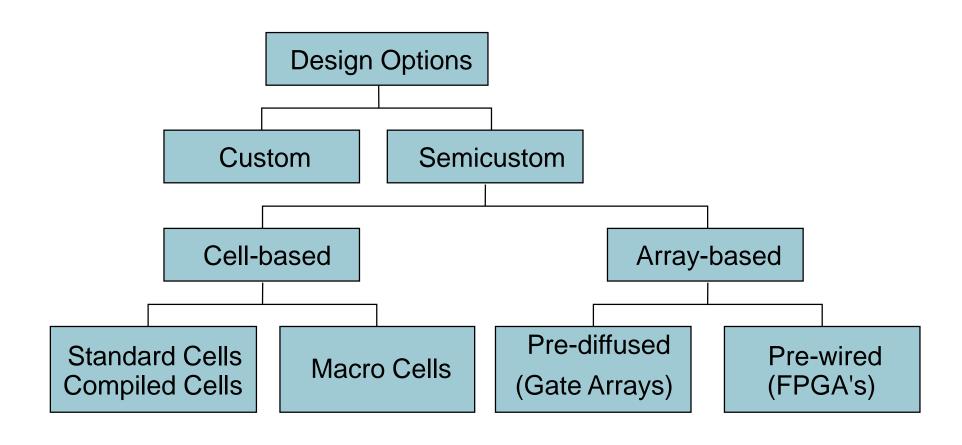
What is Reconfigurable Computing?

- Reconfigurable computing (RC) is study of SoC architectures that can adapt (after fabrication) to a specific application or application domain
 - Involves architecture, tools, CAD, design automation, algorithms, languages, etc.





Recall: Design Options



Recall: Design Approaches

	Custom	Cell-based	Pre-diffused	Pre-wired
Density	Very High	High	High	Medium
Performance	Very High	High	High	Medium
Flexibility	Very High	High	Medium	Low
Design time	Very High	Short	Short	Very Short
Manufacturing time	Medium	Medium	Short	Very Short
Cost (low volume)	Very High	High	High	Low
Cost (high volume)	Low	Low	Low	High

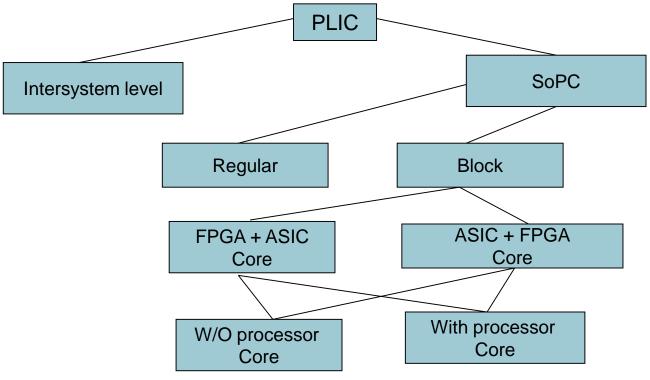
Why is RC important?

- Performance
 - Often orders of magnitude faster than microprocessors
- Low power consumption
 - A few RC devices can provide similar performance as large cluster at a fraction of the power
 - Also smaller, cheaper, etc.
- Motivating example: Novo-G
 - FPGA-based supercomputer
 - 100s of Altera/Intel FPGAs
 - 10s of Linux nodes
 - Speedups of 100,000x to 550,000x for computation biology apps (compared to 2.4 GHz Opteron)
 - Performance similar to top supercomputers
 - However, power consumption is only 8 kilowatts compared to 2-7 megawatts



RC Classification

According to Integration Level



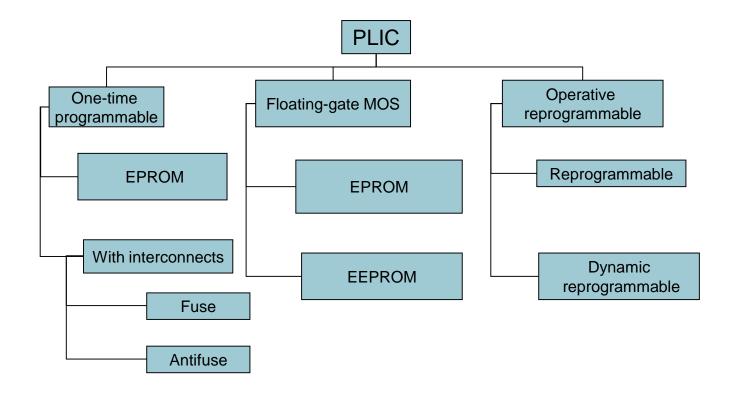
Here PLIC – Programmable Logic IC

SoPC – System-on-a-Programmable Chip



RC Classification

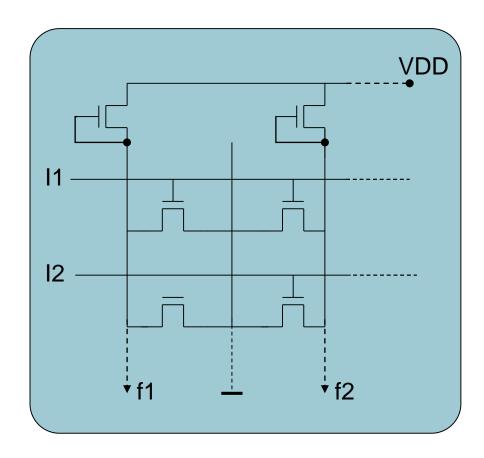
According to Reprogrammable Capacity



Implementation of Programmable Interconnects (1)

Mask Programming Interconnect

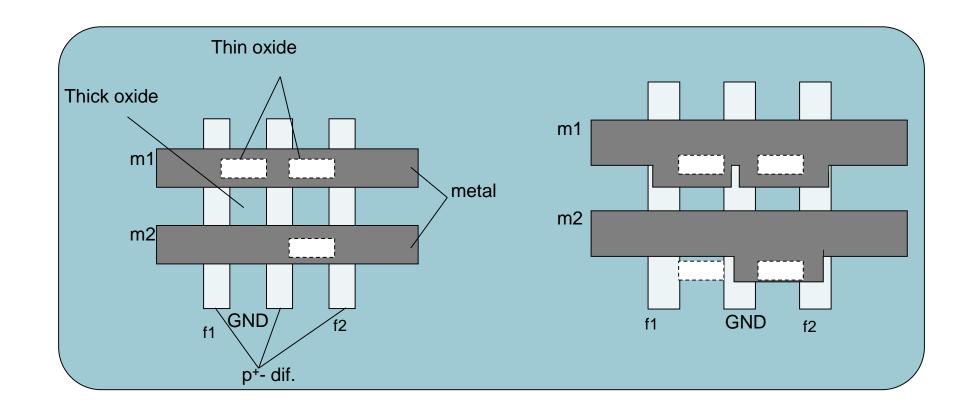
- Oxide thickness change (oxide mask changes)
- Bus width change (oxide mask changes)



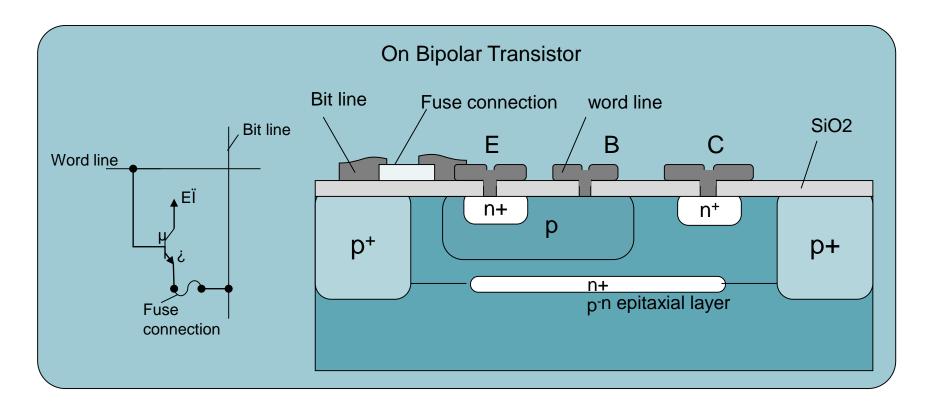
Implementation of Programmable Interconnects (2)

Oxide Thickness Change

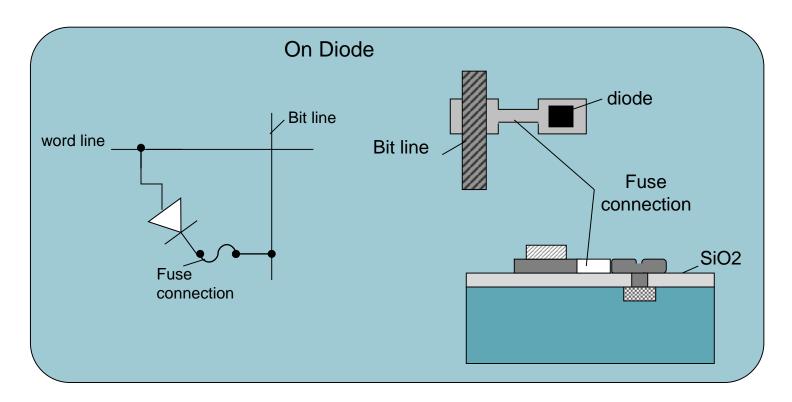
Line Width Change



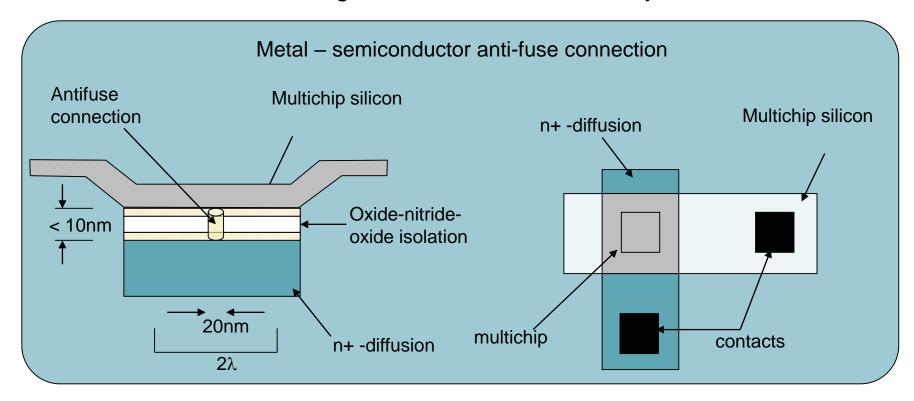
Implementation of Programmable Interconnects (3)



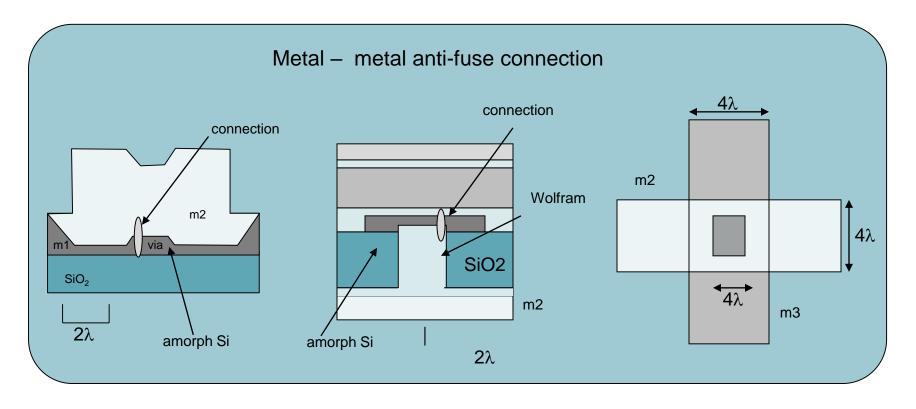
Implementation of Programmable Interconnects (4)



Implementation of Programmable Interconnects (5)

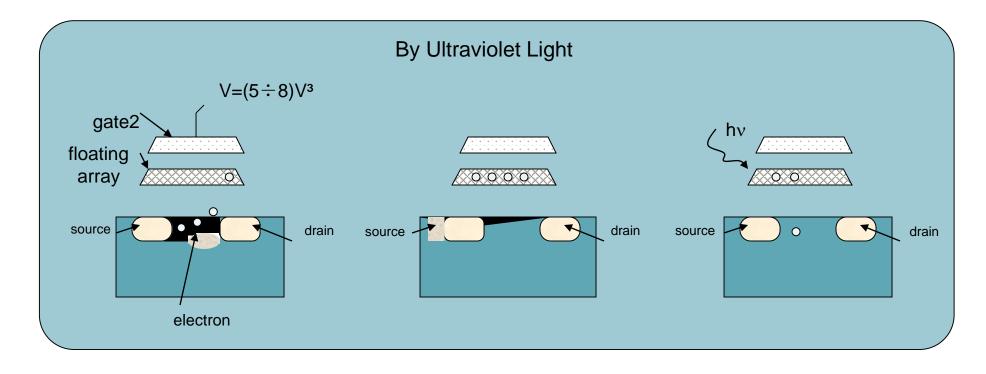


Implementation of Programmable Interconnects (6)



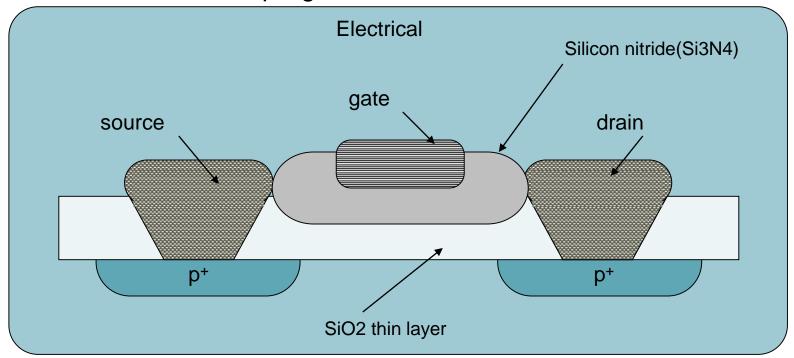
Implementation of Programmable Interconnects (7)

Reprogrammable Interconnect



Implementation of Programmable Interconnects (8)

Reprogrammable Interconnect



Zynq Reconfigurable SoC

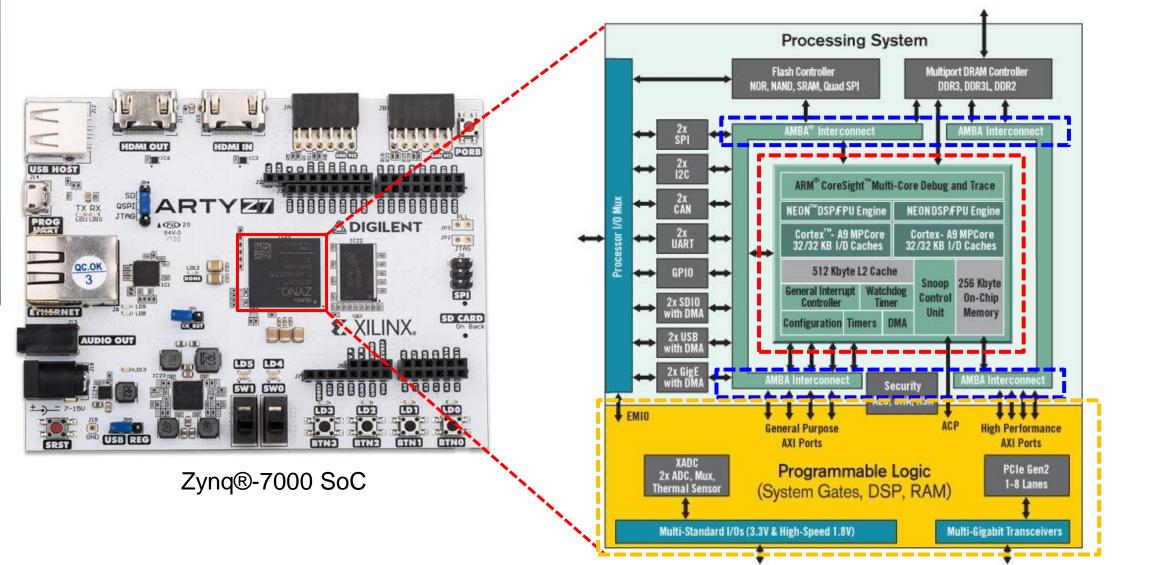
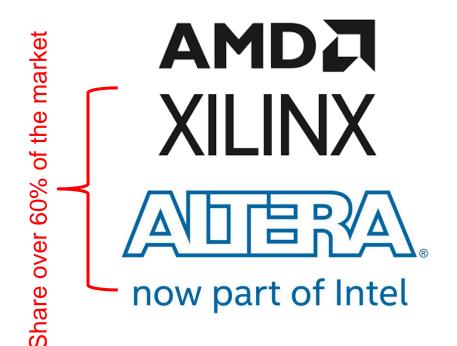




Image: Xilinx

Major Players



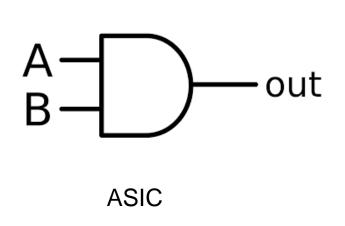


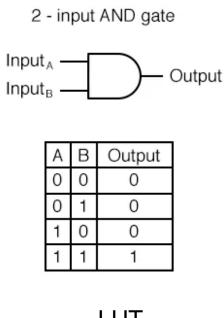


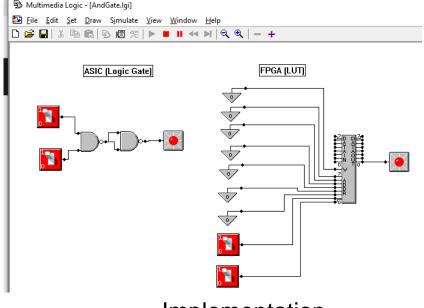


FPGA Architectures

- How can we implement any circuit in an FPGA?
 - First, focus on combinational logic
 - Example: AND gate
 - Combinational logic represented by truth table
 - What kind of hardware can implement a truth table?

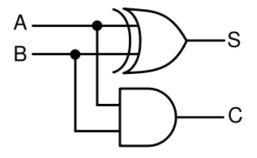






FPGA Architectures

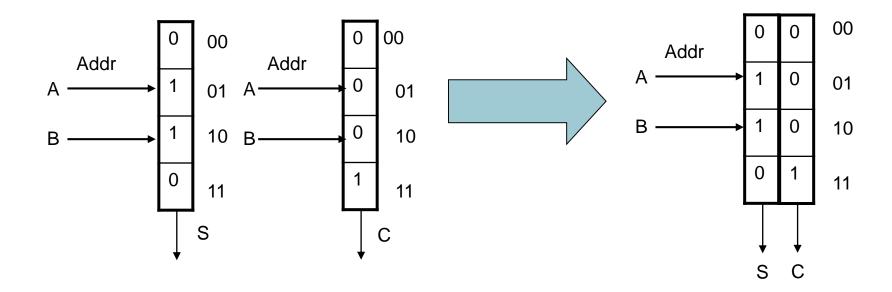
- Example: Half adder
 - Combinational logic represented by truth table
 - What kind of hardware can implement a truth table?



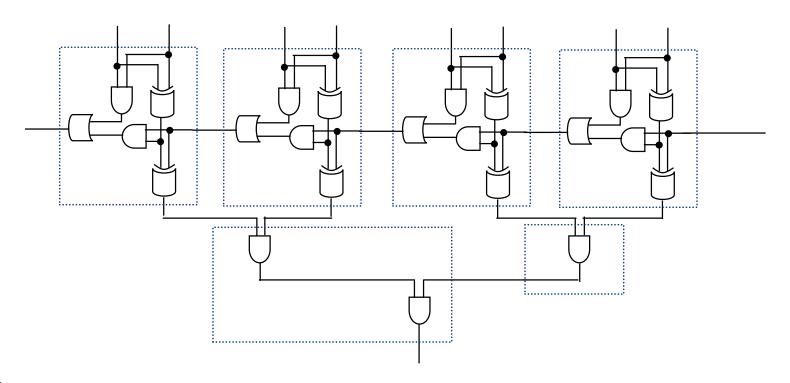
Inp	Out	
Α	В	S
0	0	0
0	1	1
1	0	1
1	1	0

Inp	Out	
Α	В	C
0	0	0
0	1	0
1	0	0
1	1	1

- Alternatively, could have used a 2-input, 2-output LUT
 - Outputs commonly use same inputs



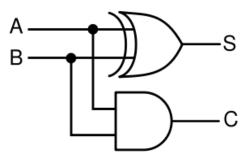
- Fortunately, we can map circuits onto multiple LUTs
 - Divide circuit into smaller circuits that fit in LUTs (same # of inputs and outputs)
 - Example: 3-input, 2-output LUTs





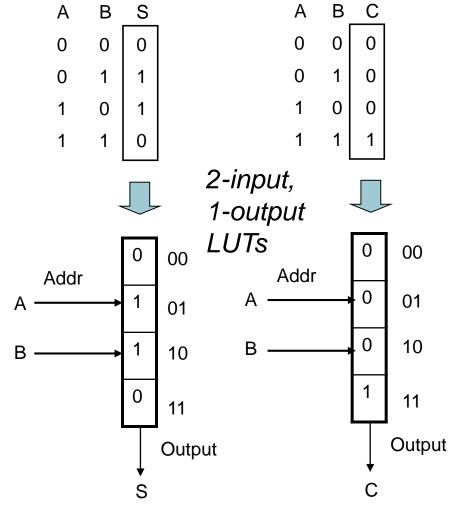
- What if circuit doesn't map perfectly?
 - More inputs in LUT than in circuit
 - Truth table handles this problem
 - Unused inputs are ignored
 - More outputs in LUT than in circuit
 - Extra outputs simply not used
 - Space is wasted, so should use multiple outputs whenever possible

- Implement truth table in small memories (LUTs)
 - Usually SRAM



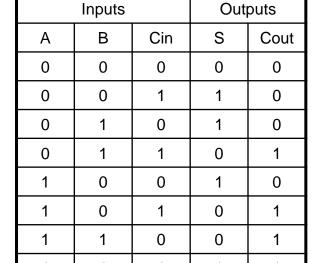
Logic inputs connect to address inputs, logic output is

memory output

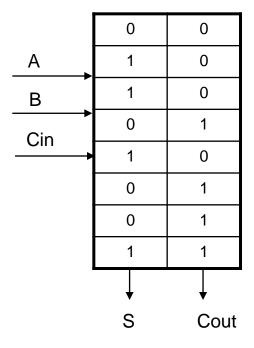


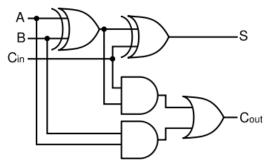
- Slightly bigger example: Full adder
 - Combinational logic can be implemented in a LUT with same number of inputs and outputs
 - 3-input, 2-ouput LUT

Truth Table



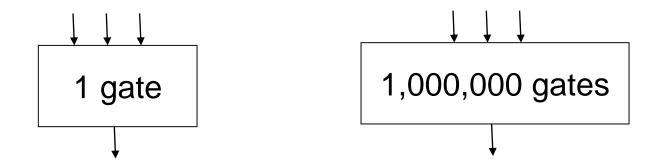
3-input, 2-output LUT





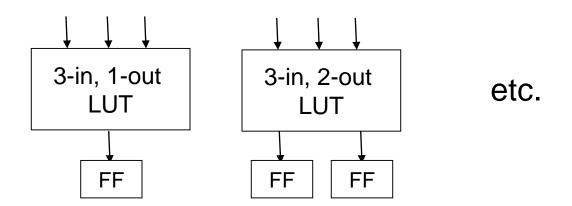
- Why aren't FPGAs just a big LUT?
 - Size of truth table grows exponentially based on # of inputs
 - 3 inputs = 8 rows, 4 inputs = 16 rows, 5 inputs = 32 rows, etc.
 - Same number of rows in truth table and LUT
 - LUTs grow exponentially based on # of inputs
- Number of SRAM bits in a LUT = 2ⁱ * o
 - i = # of inputs, o = # of outputs
 - Example: 64 input combinational logic with 1 output would require 2⁶⁴ SRAM bits
 - 1.84 x 10¹⁹
- Clearly, not feasible to use large LUTs
 - So, how do FPGAs implement logic with many inputs?

- Important Point
 - The number of gates in a circuit has no effect on the mapping into a LUT
 - All that matters is the number of inputs and outputs
 - Unfortunately, it isn't common to see large circuits with a few inputs

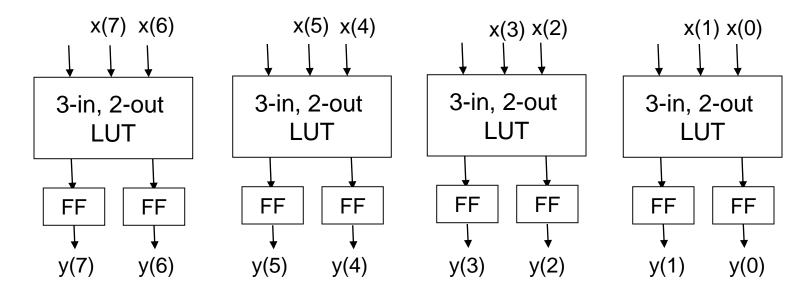


Both of these circuits can be implemented in a single 3-input, 1-output LUT

- Problem: How to handle sequential logic
 - Truth tables don't work
- Possible solution:
 - Add a flip-flop to the output of LUT

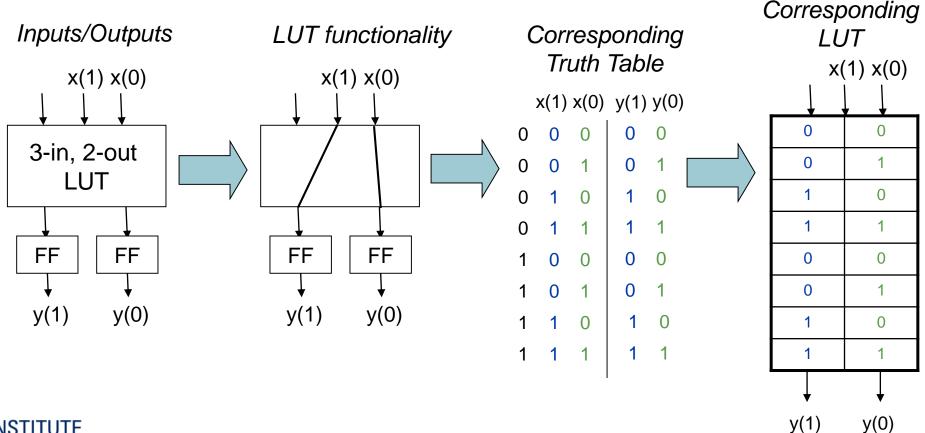


- Example: 8-bit register using 3-input, 2-output LUTs
 - Input: x, Output: y

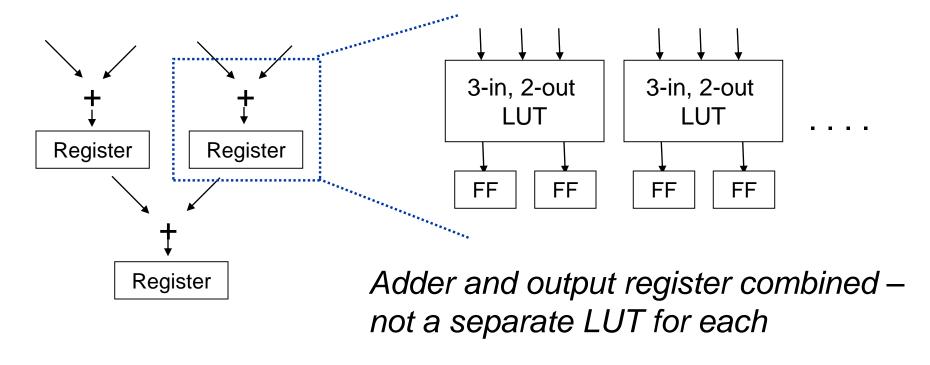


What does LUT need to do to implement register?

- Example, cont.
 - LUT simply passes inputs to appropriate output



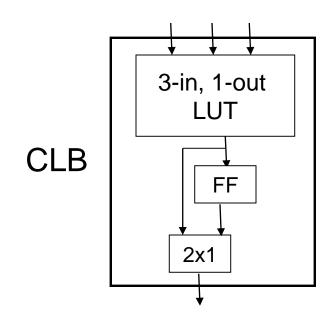
- Isn't it a waste to use LUTs for registers?
- YES! (when it can be used for something else)
 - Commonly used for pipelined circuits
 - Example: Pipelined adder



- Existing FPGAs don't have a flip flop connected to LUT outputs
- Why not?
 - Flip flop has to be used!
 - Impossible to have pure combinational logic
 - Adds latency to circuit
- Actual Solution:
 - Configurable Logic Blocks (CLBs)

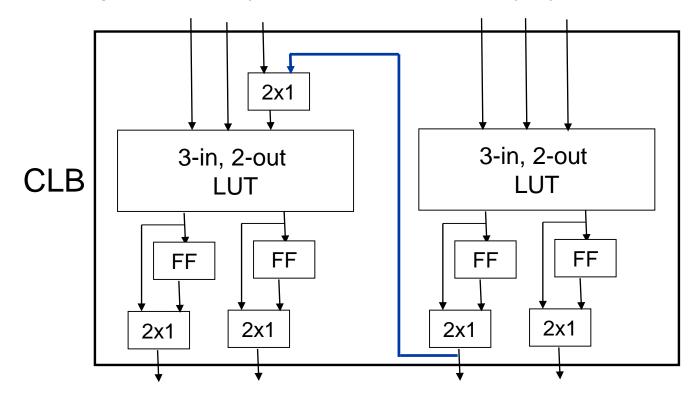


- CLBs: the basic FPGA functional unit
 - First issue: How to make flip-flop optional?
 - Simplest way: use a mux
 - Circuit can now use output from LUT or from FF
 - Where does select come from? (will be answered shortly)

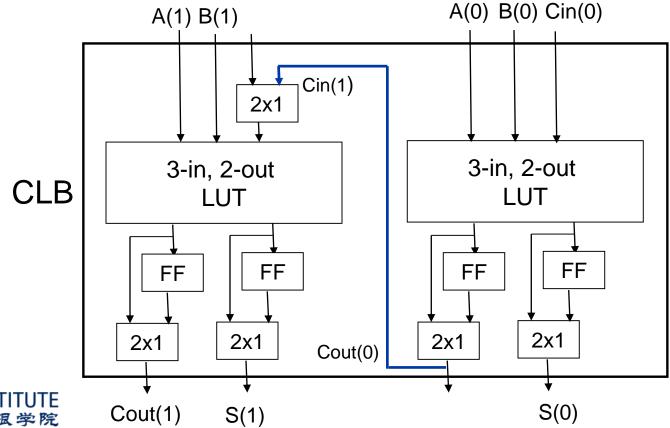




- CLBs usually contain more than 1 LUT
 - Why?
 - Efficient way of handling common I/O between adjacent LUTs
 - Saves routing resources (we haven't discussed yet)



- Example: Ripple-carry adder
 - Each LUT implements 1 full adder
 - Use efficient connections between LUTs for carry signals



- CLBs often have specialized connections between adjacent CLBs
 - Further improves carry chains
 - Avoids routing resources
- Some commercial CLBs even more complex
 - Xilinx Virtex 4 CLB consists of 4 "slices"
 - 1 slice = 2 LUTs + 2 FFs + other stuff
 - 1 Virtex 4 CLB = 8 LUTs
 - Altera devices has LABs (Logic Array Blocks)
 - Consist of 16 LEs (logic elements) which each have 4 input LUTs

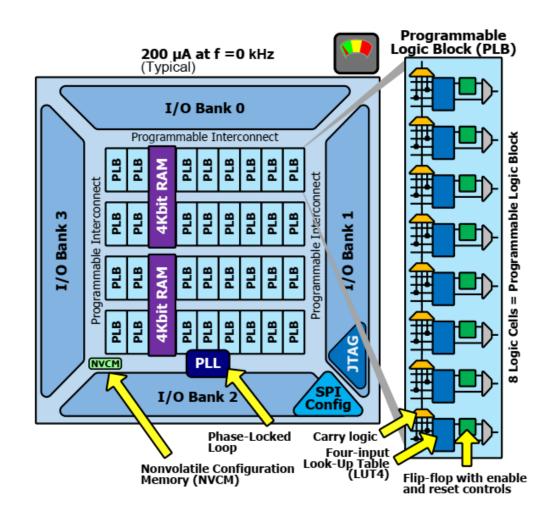
CLB Examples

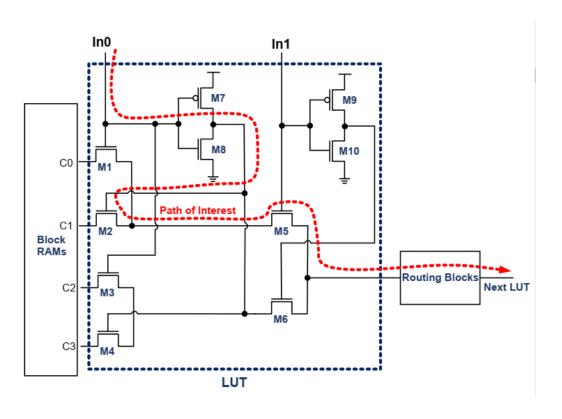
- Virtex 4 CLB
 - http://www.xilinx.com/support/documentation/user_guides/ug070.pdf (pg. 183)
- Altera Stratix 5
 - http://www.altera.com/literature/hb/stratix-v/stratix5_handbook.pdf (pg. 10)
- Virtex 7 CLB
 - http://www.xilinx.com/support/documentation/user_guides/ug474_7Series_
 CLB.pdf (pg. 13)
 - http://www.xilinx.com/csi/training/7_series_CLB_architecture.htm

CLB Examples

- Intel/Altera Arria 10 CLB (ALM)
 - https://www.intel.com/content/dam/www/programmable/us/en/pdfs/literatur e/hb/arria-10/a10_overview.pdf (pg. 18)
- Intel/Altera Stratix 10
 - https://www.intel.com/content/dam/www/programmable/us/en/pdfs/literatur e/hb/stratix-10/s10-overview.pdf (pg. 24)
- Xilinx UltraSacle
 - https://www.xilinx.com/support/documentation/data_sheets/ds890ultrascale-overview.pdf (pg.36)

An example





2input, 1output LUT



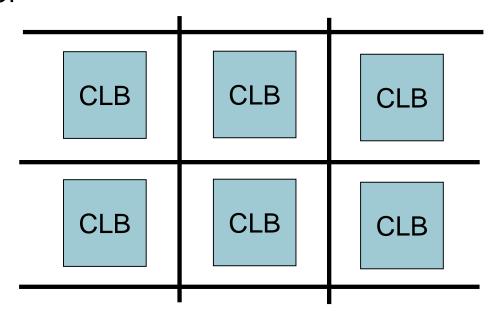
What Else?

- Basic building block is CLB
 - Can implement combinational+sequential logic
 - All circuits consist of combinational and sequential logic
- So what else is needed?

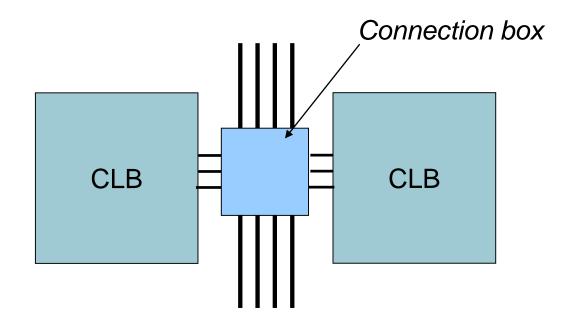
- FPGAs need some way of connecting CLBs together
 - Reconfigurable interconnect
 - But, we can only put fixed wires on a chip
- Problem: How to make reconfigurable connections with fixed wires?
 - Main challenge:
 - Should be flexible enough to support almost any circuit



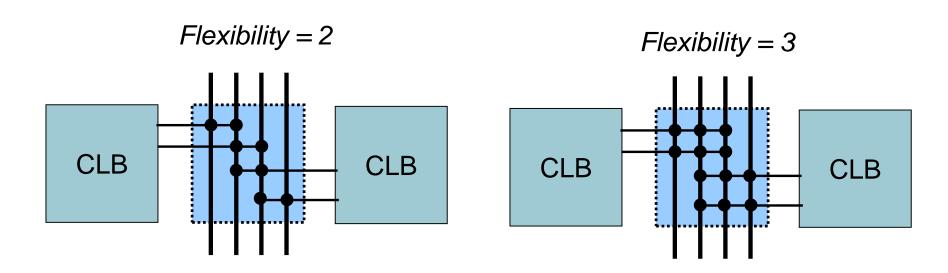
- Problem 2: If FPGA doesn't know which CLBs will be connected, where does it put wires?
- Solution:
 - Put wires everywhere!
 - Referred to as channel wires, routing channels, routing tracks, many others
 - CLBs typically arranged in a grid, with wires on all sides
 - Remind of NoC!



- Problem 3: How to connect CLB to wires?
- Solution: Connection box
 - Device that allows inputs and outputs of CLB to connect to different wires

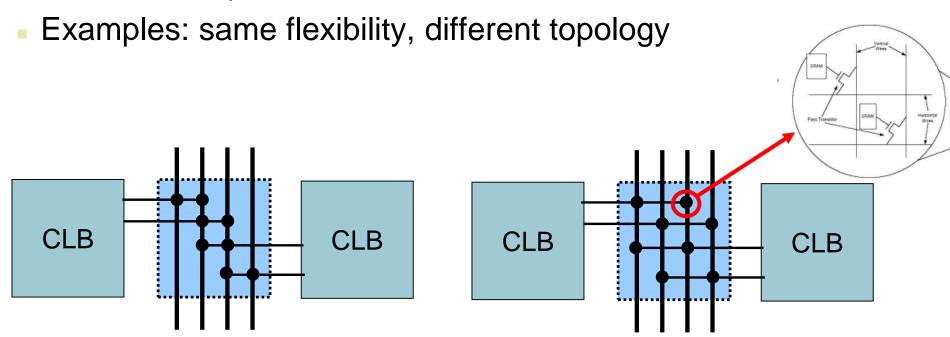


- Connection box characteristics
 - Flexibility
 - The number of wires a CLB input/output can connect to





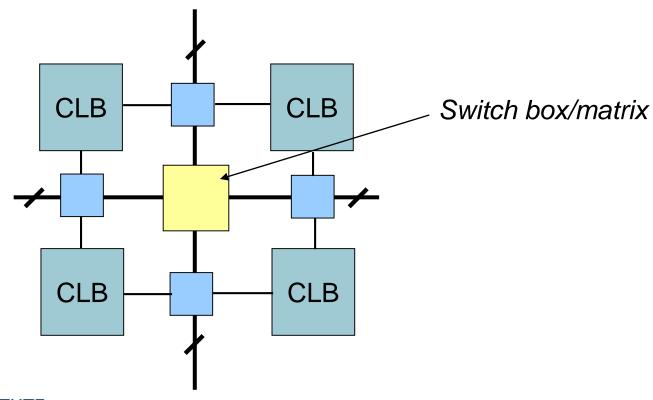
- Connection box characteristics
 - Topology
 - Defines the specific wires each CLB I/O can connect to





- Connection boxes allow CLBs to connect to routing wires
 - But, that only allows us to move signals along a single wire
 - Not very useful
- Problem 4: How do FPGAs connect wires together?

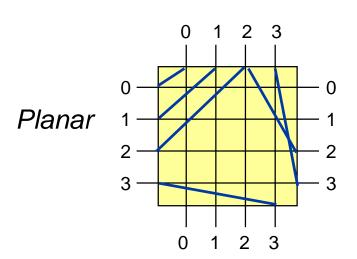
- Solution: Switch boxes, switch matrices
 - Connects horizontal and vertical routing channels

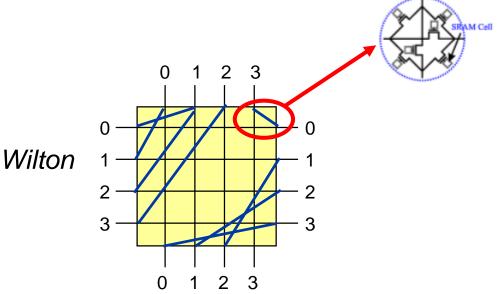




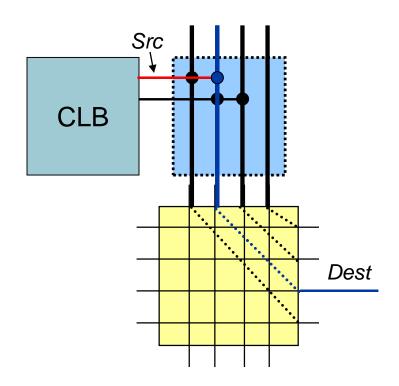
- Switch boxes
 - Flexibility defines how many wires a single wire can connect to
 - Topology defines which wires can be connected
 - Planar/subset switch box: only connects tracks with same id/offset (e.g. 0 to 0, 1 to 1, etc.)

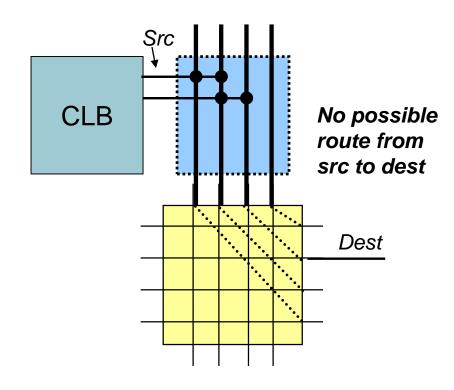
Wilton switch box: connects tracks with different offsets



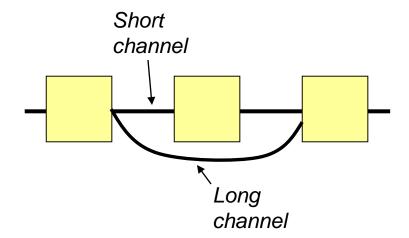


- Why do flexiblity and topology matter?
 - Routability: a measure of the number of circuits that can be routed
 - Higher flexibility = better routability
 - Wilton switch box topology = better routability





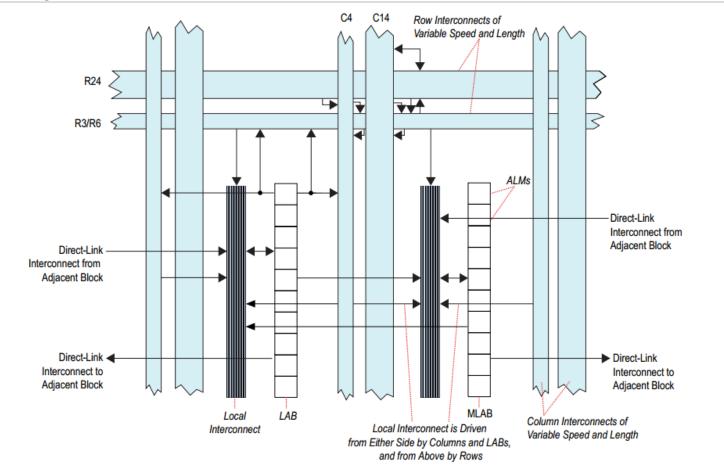
- Switch boxes
 - Short channels
 - Useful for connecting adjacent CLBs
 - Long channels
 - Useful for connecting CLBs that are separated
 - Allows for reduced routing delay for non-adjacent CLBs



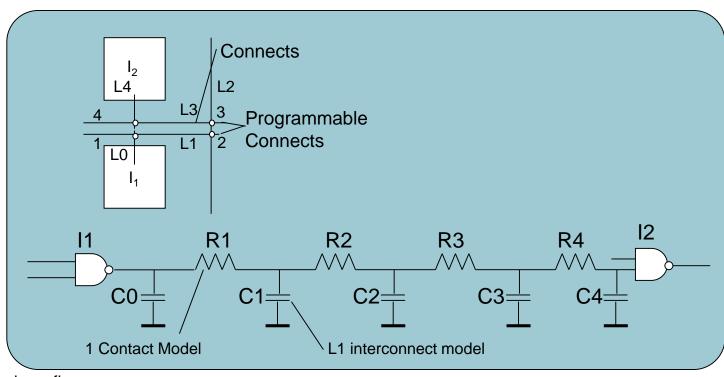
Interconnect Example

- Altera provides long tracks of length 3, 4, 6, 14, 24 along with local interconnect (short tracks)
- Image from Stratix V handbook. LAB = CLB, ALM = LUT

This figure shows an overview of the Stratix V LAB and MLAB structure with the LAB interconnects.



Delays in Reconfigurable Interconnects (1)



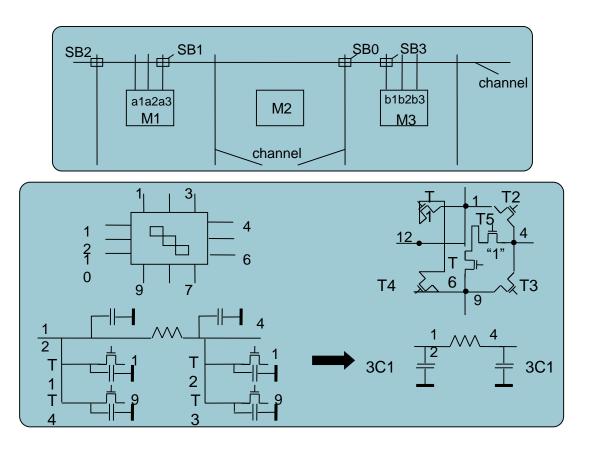
From the above figure:

 $\tau = R1 \cdot C1 + (R1 + R2) \cdot C2 + (R1 + R2 + R3) \cdot C3 + (R1 + R2 + R3 + R4) \cdot C4$

If R=R1=R2=R3=R4 & C=C1=C2=C3=C4, then τ =R · C+2R · C+3R · C+4R · C=10R · C

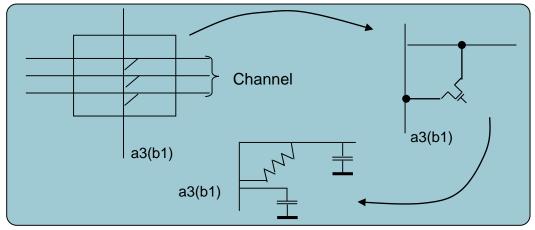


Delays in Reconfigurable Interconnects (2)

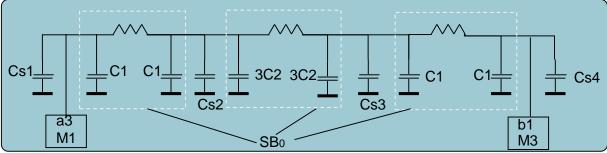


Delays in Reconfigurable Interconnects (3)





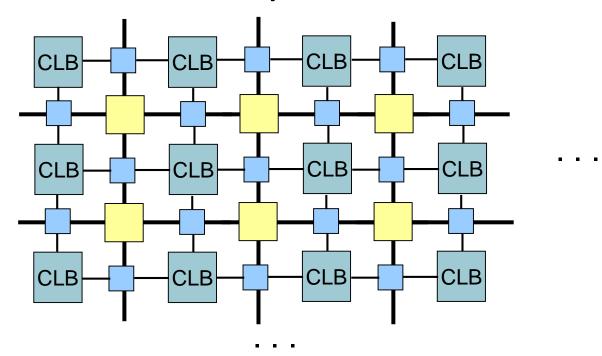
General equivalent circuit



From the above figure: τ =R5C+2R5C+3R2C=21RC

FPGA Fabrics

- FPGA layout called a "fabric"
 - 2-dimensional array of CLBs and programmable interconnect
 - Sometimes referred to as an "island style" architecture



- Can implement any circuit
 - But, should fabric include something else?



FPGA Fabrics

- What about memory?
 - Could use FF's in CLBs to create a memory
 - Example: Create a 1 MB memory with:
 - CLB with a single 3-input, 2-output LUT
 - Each CLB = 2 bits of memory (because of 2 outputs)
 - Total CLBs = (1 MB * 8 bits/byte) / 2 bits/CLB
 - 4 million CLBs!!!!
 - FPGAs commonly have tens of thousands of LUTs
 - Large devices have 100-200k LUTs
 - State-of-the-art devices ~800k LUTs
 - Even if FPGAs were large enough, using a chip to implement 1 MB of memory is not smart
 - Conclusion:
 - Bad Idea!! Huge waste of resources!

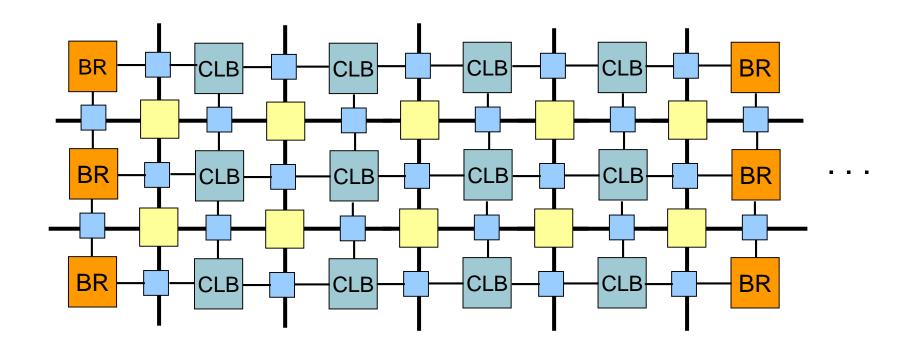
FPGA Memory Components

- Solution 1: Use LUTs for logic or memory
 - LUTs are small SRAMs, why not use them as memory?
 - Xilinx refers to as distributed RAM
 - Intel/Altera refers to as MLABs
- Solution 2: Include dedicated RAM components in the FPGA fabric
 - Xilinx refers to as Block RAM
 - Can be single/dual-ported
 - Can be combined into arbitrary sizes
 - Can be used as FIFO
 - Different clock speeds for reads/writes
 - Altera has Memory Blocks
 - M4K: 4k bits of RAM
 - Others: M9K, M20k, M144K



FPGA Memory Components

- Fabric with Block RAM
 - Block RAM can be placed anywhere
 - Typically, placed in columns of the fabric



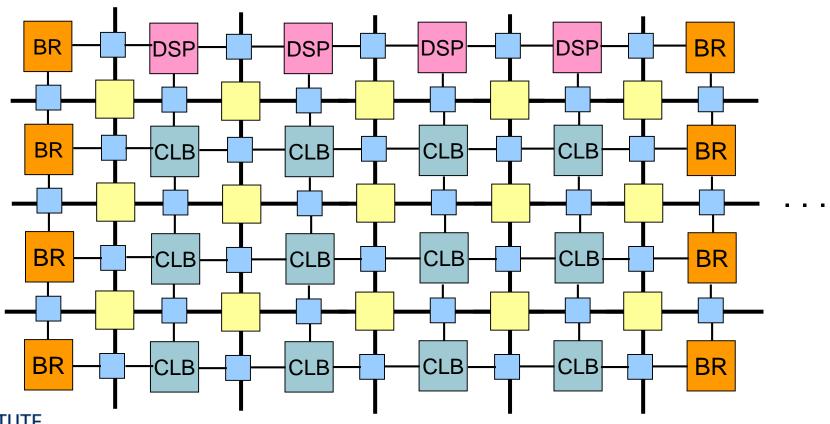


DSP Components

- FPGAs commonly used for DSP apps
 - Makes sense to include custom DSP units instead of mapping onto LUTs
 - Custom unit = faster/smaller
- Example: Xilinx DSP48
 - Includes multipliers, adders, subtractors, etc.
 - 18x18 multiplication
 - 48-bit addition/subtraction
 - Provides efficient way of implementing
 - Add/subtract/multiply
 - MAC (Multiply-accumulate)
 - Barrel shifter
 - FIR Filter
 - Square root
 - Etc.
- Older Altera devices have multiplier blocks
 - Can be configured as 18x18 or 2 separate 9x9 multipliers
- Intel Arria/Stratix 10 have floating-point DSPs
 - https://www.intel.com/content/dam/www/programmable/us/en/pdfs/literature/ug/ug-a10dsp.pdf

Example Fabric

- Existing FPGAs are 2-dimensional arrays of CLBs, DSP, Block RAM, and programmable interconnect
 - Actual layout/placement differs for different FPGAs
 - Specialized resources tend to be in columns not rows



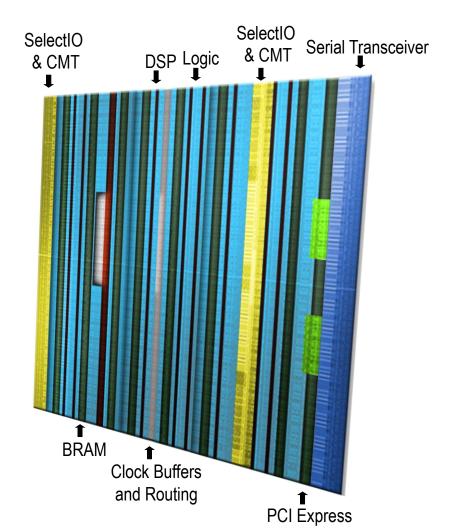


Other resources

- I/O
 - Virtex 7 has 1,200 pins
 - Communication is still often a bottleneck
 - Pins don't increase with new FPGAs, but logic does
 - Trend: High-speed serial transceivers
- Clock resources
 - Using reconfigurable interconnect for clock introduces timing problems
 - Skew, jitter
 - FPGAs often provided clock trees, both globally and locally
 - e.g. Virtex 7
 http://www.xilinx.com/support/documentation/user_guides/ug472_7Series_Clocking.pdf

Example Fabrics

Virtex 7 (image from Xilinx 7-series overview)

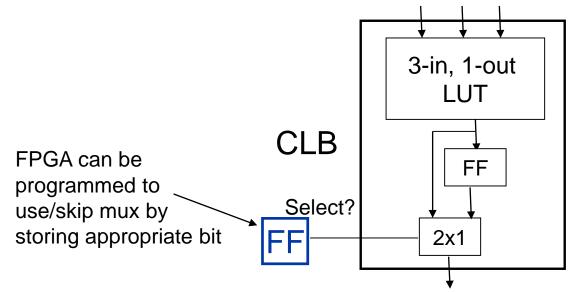


CMT: clock management tiles

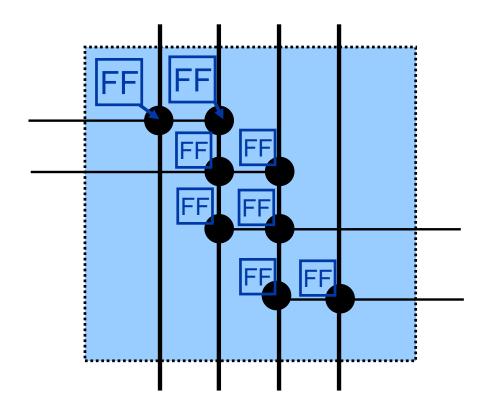
- How to program/configure FPGA to implement circuit?
 - So far, we've mapped a circuit onto FPGA fabric
 - Known as technology mapping
 - Process of converting a circuit in one representation into a representation that corresponds to physical components
 - Gates to LUTs
 - Memory to Block RAMs
 - Multiplications to DSP48s
 - Etc.
 - But, we need some way of configuring each component to behave as desired
 - Examples:
 - How to store truth tables in LUTs?
 - How to connect wires in switch boxes?
 - Etc.



- General Idea: include FF's in fabric to control programmable components
 - Example: CLB
 - Need a way to specify select for mux



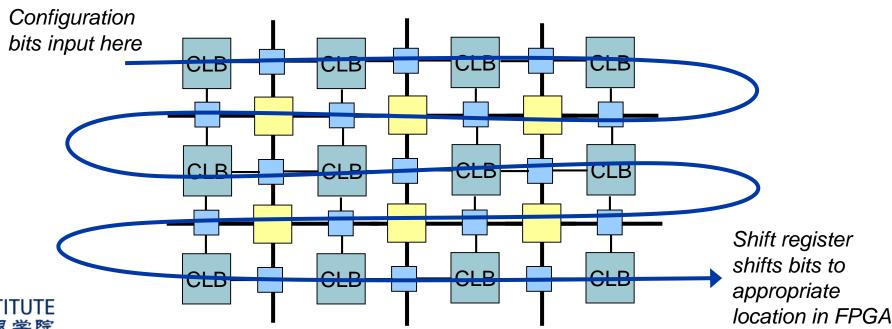
- Example 2:
 - Connection/switch boxes
 - Need FFs to specify connections



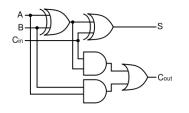
- FPGAs programmed with a "bitfile"
 - File containing all information needed to program FPGA
 - Contains bits for each control FF
 - Also, contains bits to fill LUTs
- But, how do you get the bitfile into the FPGA?
 - > 10k LUTs
 - Small number of pins



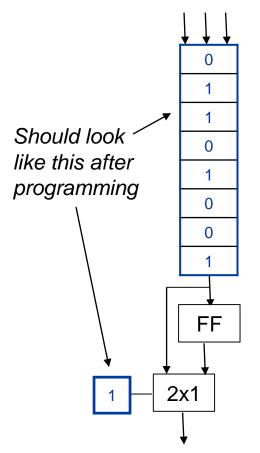
- Solution: Shift Registers
 - General Idea
 - Make a huge shift register out of all programmable components (LUTs, control FFs)
 - Shift in bitfile one bit at a time

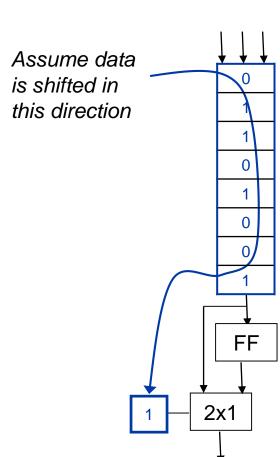


- Example:
 - Program CLB with 3-input, 1-output LUT to implement sum output of full adder

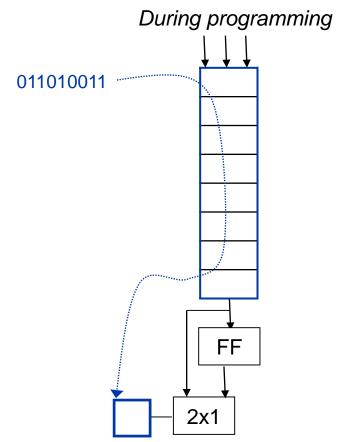


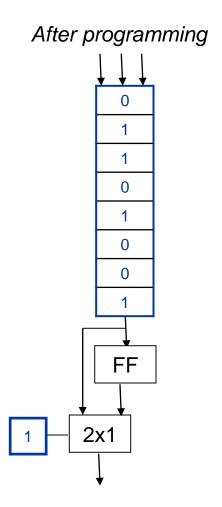
In			Out
Α	В	Cin	S
0	0	0	0
0	0	1	1
0	1	0	1
0	1	1	0
1	0	0	1
1	0	1	0
1	1	0	0
1	1	1	1



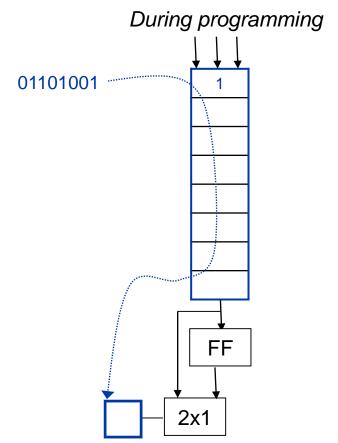


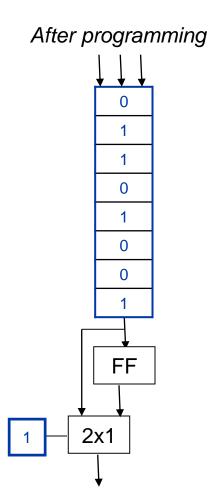
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register





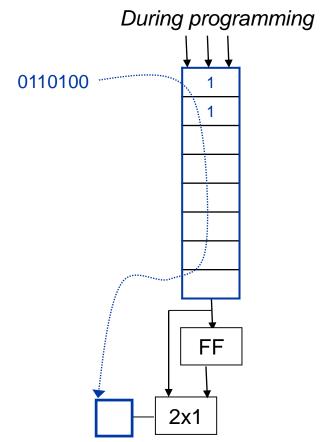
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register

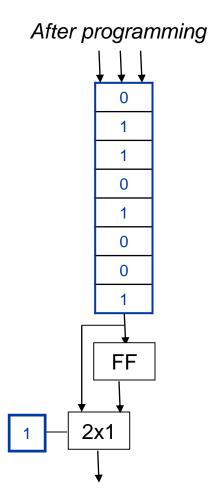




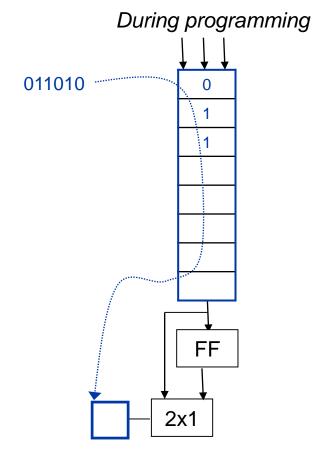


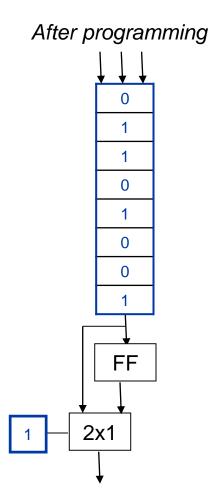
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register





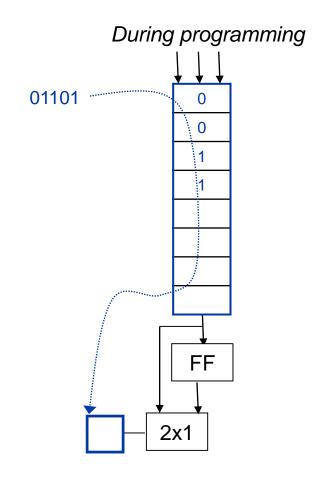
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register

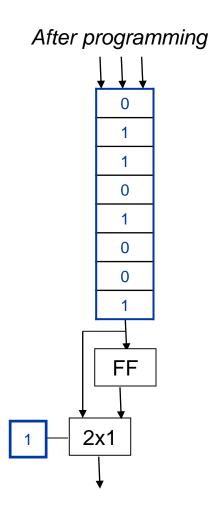




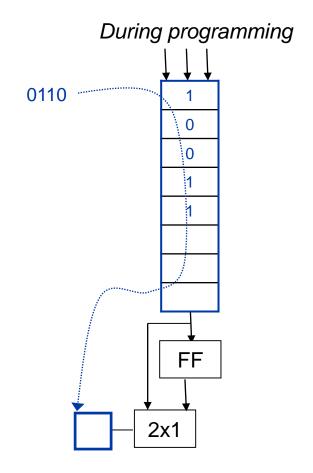


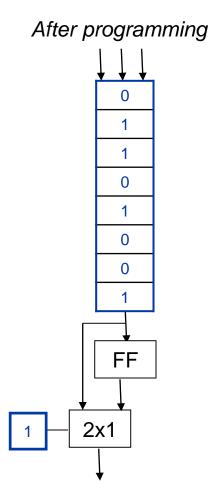
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register





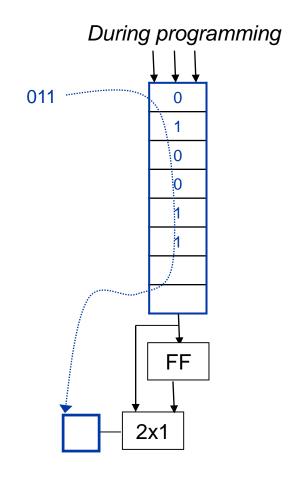
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register

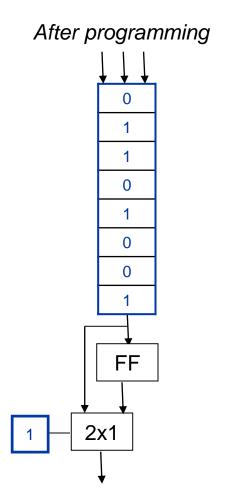




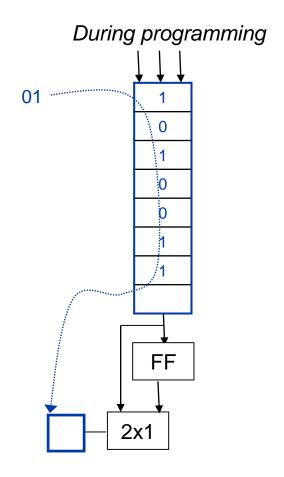


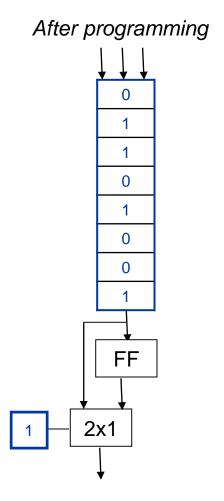
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register



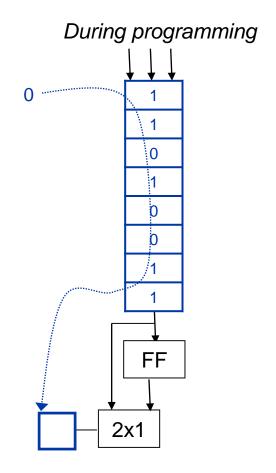


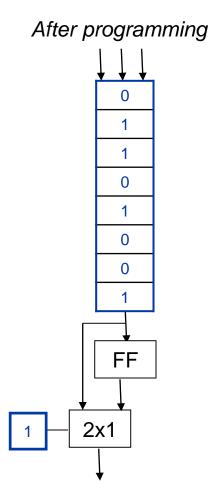
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register





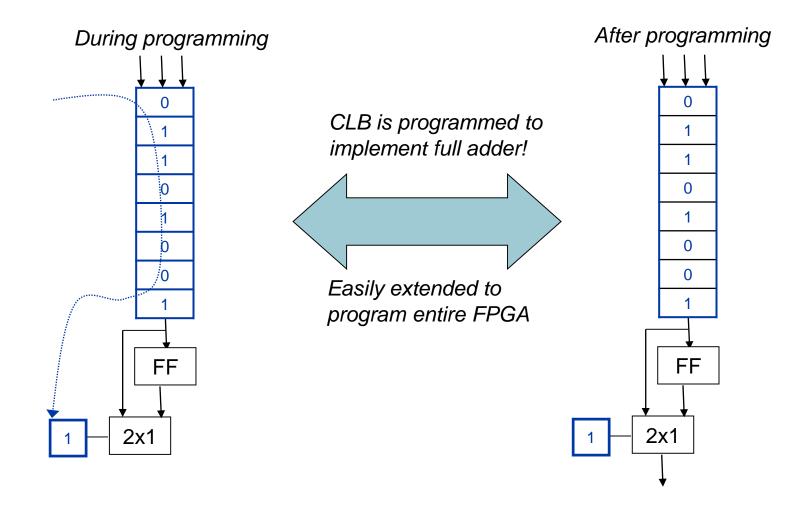
- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register







- Example, Cont:
 - Bitfile is just a sequence of bits based on order of shift register



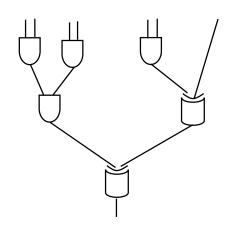
- Problem: Reconfiguring FPGA is slow
 - Shifting in 1 bit at a time not efficient
 - Bitfiles can be greater than 1 MB
 - Eliminates one of the main advantages of RC
 - Partial reconfiguration
 - With shift registers, entire FPGA has to be reconfigured
- Solutions?
 - Virtex II allowed columns to be reconfigured
 - Virtex IV allowed custom regions to be reconfigured
 - Requires a lot of user effort
 - Better tools needed



- LUTs with many inputs can implement large circuits efficiently
 - Why not just use LUTs with many inputs?
- High flexibility in routing resources improves routability
 - Why not just allow all possible connections?
- Answer: architectural tradeoffs
 - Anytime one component is increased/improved, there is less area for other components
 - Larger LUTs => less total LUTs, less routing resources
 - More Block RAM => less LUTs, less DSPs
 - More DSPs => less LUTs, less Block RAM
 - Etc.

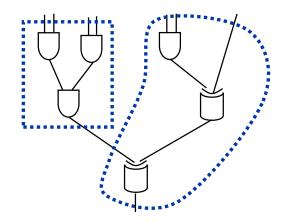


- Example:
 - Determine best LUTs for following circuit
 - Choices
 - 4-input, 2-output LUT (delay = 2 ns)
 - 5-input, 2-output LUT (delay = 3 ns)
 - Assume each SRAM cell is 6 transistors
 - 4-input LUT = 6 * 2⁴ * 2 = 192 transistors
 - 5-input LUT = $6 * 2^5 * 2 = 384$ transistors



- Example:
 - Determine best LUTs for following circuit
 - Choices
 - 4-input, 2-output LUT (delay = 2 ns)
 - 5-input, 2-output LUT (delay = 3 ns)
 - Assume each SRAM cell is 6 transistors
 - 4-input LUT = 6 * 2⁴ * 2 = 192 transistors
 - 5-input LUT = $6 * 2^5 * 2 = 384$ transistors

5-input LUT

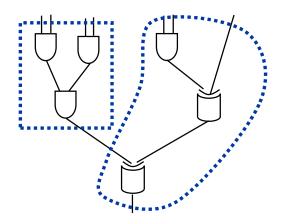


Propagation delay = 6 ns

Total transistors = 384 * 2 = 768

- Example:
 - Determine best LUTs for following circuit
 - Choices
 - 4-input, 2-output LUT (delay = 2 ns)
 - 5-input, 2-output LUT (delay = 3 ns)
 - Assume each SRAM cell is 6 transistors
 - 4-input LUT = 6 * 2⁴ * 2 = 192 transistors
 - 5-input LUT = $6 * 2^5 * 2 = 384$ transistors

4-input LUT



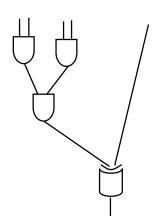
Propagation delay = 4 ns

Total transistors = 192 * 2 = 384

4-input LUTs are 1.5x faster and use 1/2 the area

Example 2

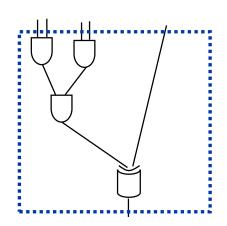
- Determine best LUTs for following circuit
 - Choices
 - 4-input, 2-output LUT (delay = 2 ns)
 - 5-input, 2-output LUT (delay = 3 ns)
 - Assume each SRAM cell is 6 transistors
 - 4-input LUT = 6 * 2⁴ * 2 = 192 transistors
 - 5-input LUT = $6 * 2^5 * 2 = 384$ transistors



Example 2

- Determine best LUTs for following circuit
 - Choices
 - 4-input, 2-output LUT (delay = 2 ns)
 - 5-input, 2-output LUT (delay = 3 ns)
 - Assume each SRAM cell is 6 transistors
 - 4-input LUT = 6 * 2⁴ * 2 = 192 transistors
 - 5-input LUT = $6 * 2^5 * 2 = 384$ transistors

5-input LUT



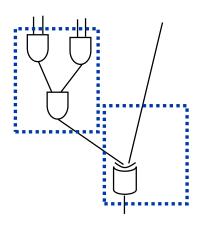
Propagation delay = 3 ns

Total transistors = 384

Example 2

- Determine best LUTs for following circuit
 - Choices
 - 4-input, 2-output LUT (delay = 2 ns)
 - 5-input, 2-output LUT (delay = 3 ns)
 - Assume each SRAM cell is 6 transistors
 - 4-input LUT = 6 * 2⁴ * 2 = 192 transistors
 - 5-input LUT = $6 * 2^5 * 2 = 384$ transistors

4-input LUT



Propagation delay = 4 ns

Total transistors = 384 transistors

5-input LUTs are 1.3x faster and use same area



- Large LUTs
 - Fast when using all inputs
 - Wastes transistors otherwise
- Must also consider total chip area
 - Wasting transistors may be ok if there are plently of LUTs
 - Virtex V uses 6 input LUTs
 - Virtex IV uses 4 input LUTs

- How to design FPGA fabric?
 - There is no overall best
 - Design fabric based on different domains
 - DSP will require many of DSP units
 - HPC may require balance of units
 - SoCs may require microprocessors

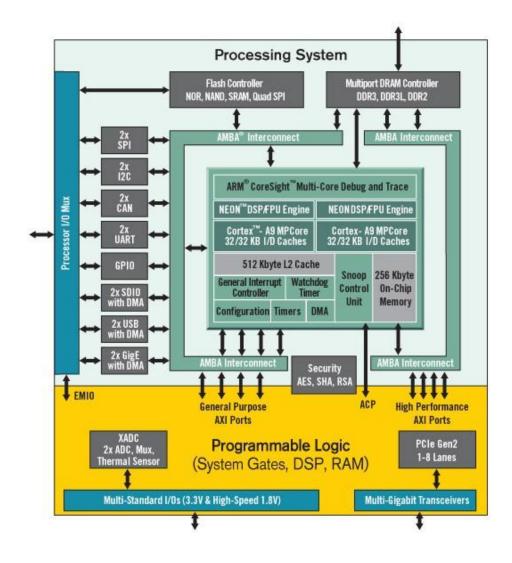
Examples:

- Xilinx Virtex IV
 - LX designed for logic intensive apps
 - SX designed for signal processing apps
 - FX designed for embedded systems apps
 - Has 450 MHz PowerPC cores embedded in fabric
- Xilinx 7 Series
 - Artix, Kintex, Virtex

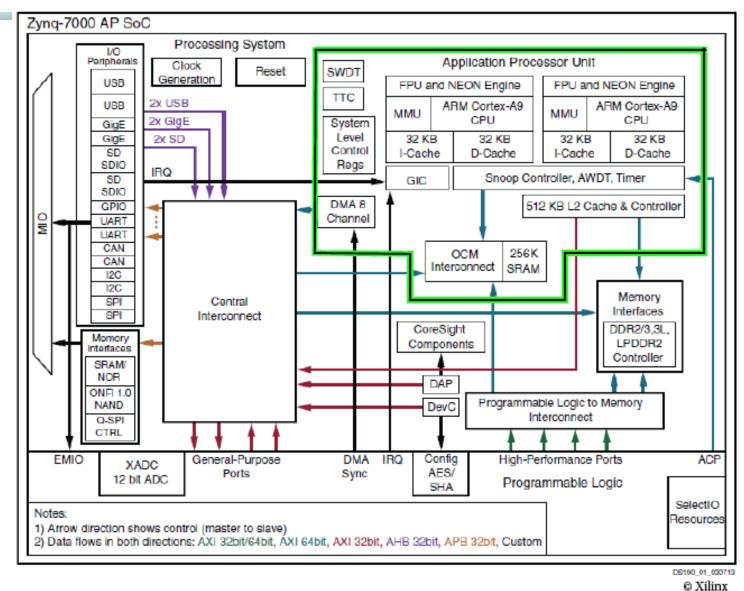


Come back to Zynq

- Combines ARM processor with programmable logic (PL)
 - Artix FPGA
 - DRAM controller
 - PCIe controller
 - Other peripherals

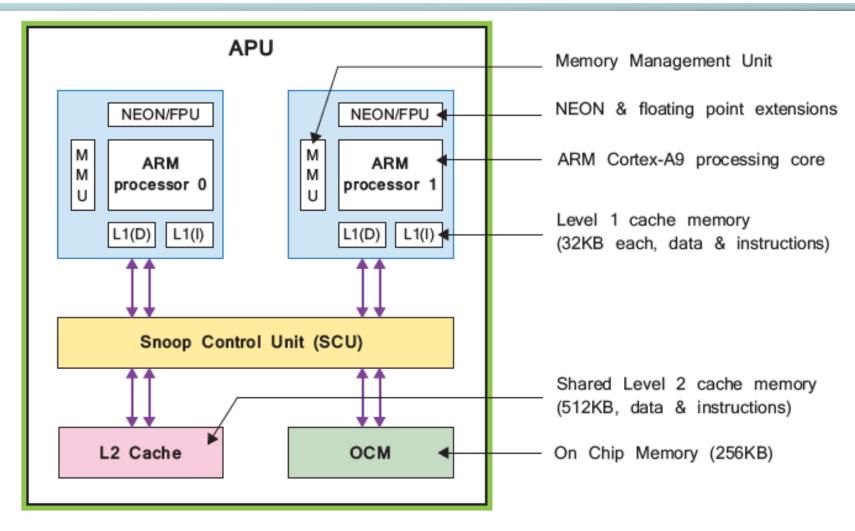


The Zynq Processing System



Source: The Zyng Book

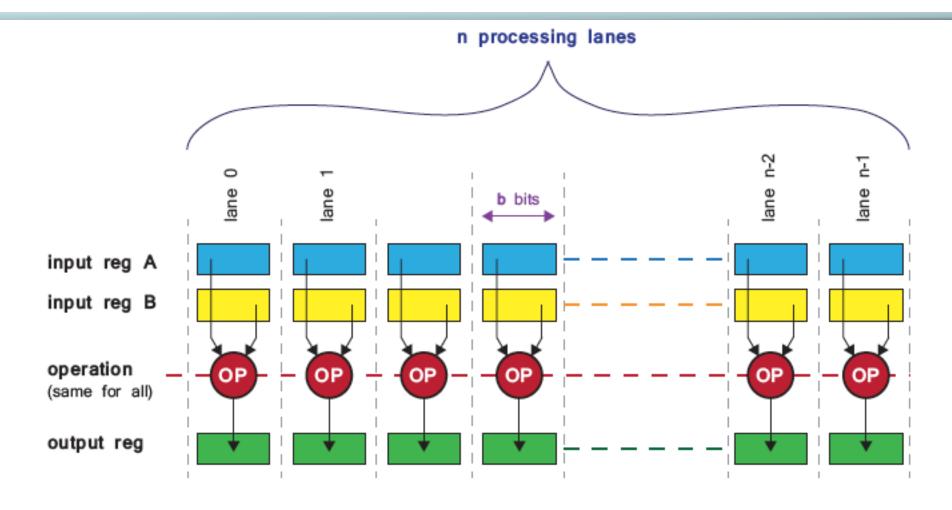
Application Processing Unit (APU)



APU programming is through Xlinx SDK



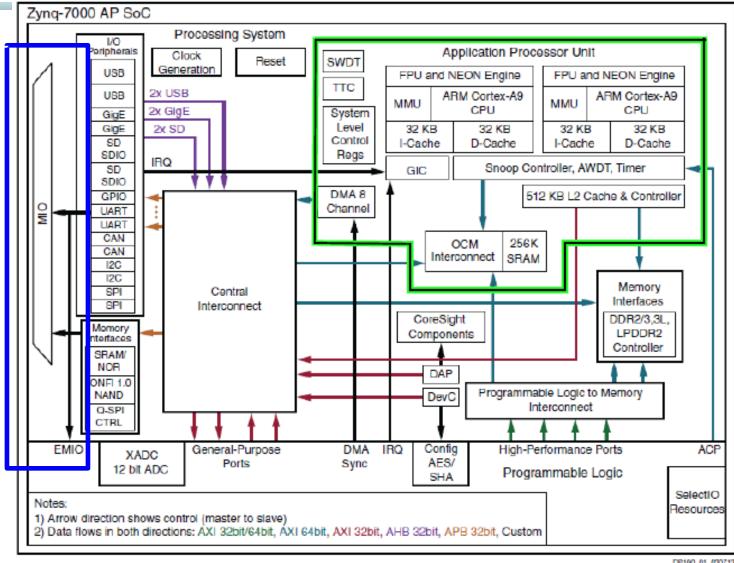
NEON Co-Processor



SIMD Execution for media and DSP applications



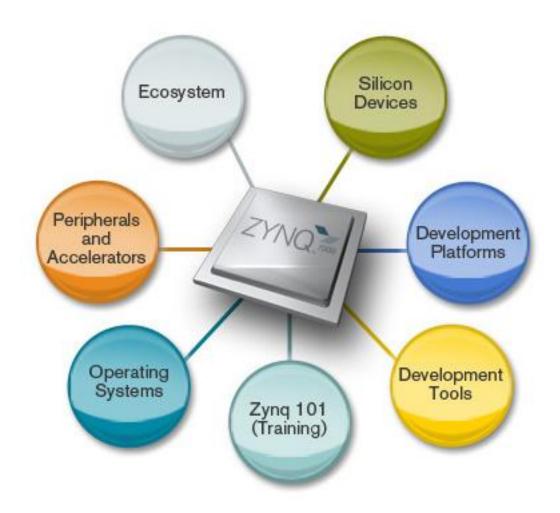
PS External Interfaces: MIO



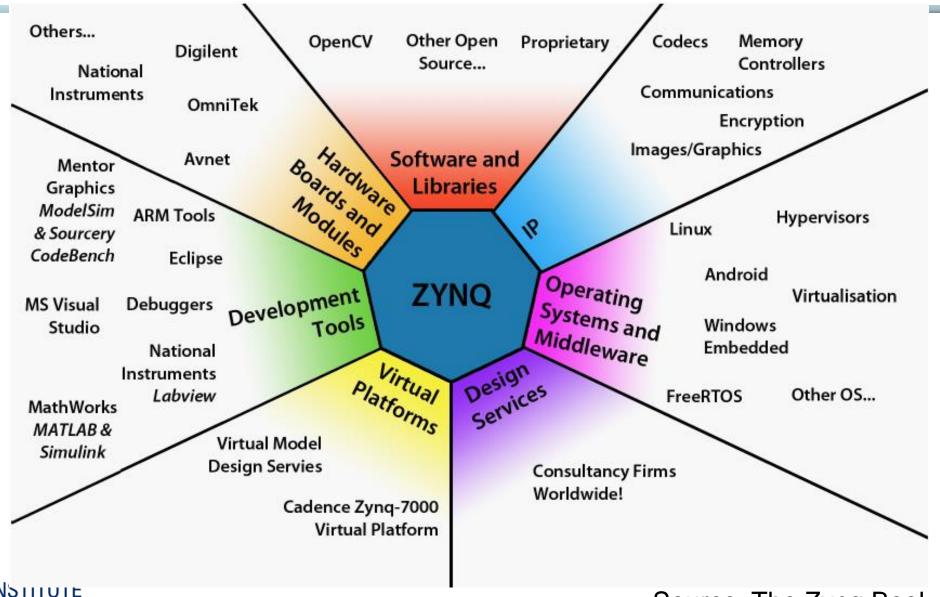


© Xilinx

Zynq SoC Ecosystem



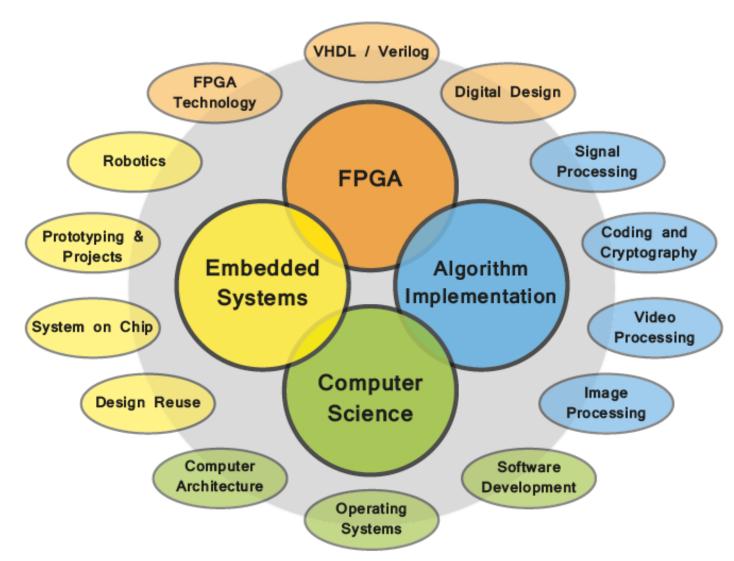
Zynq SoC Ecosystem





Source: The Zynq Book

Academic Subjects to which Zynq is Relevant

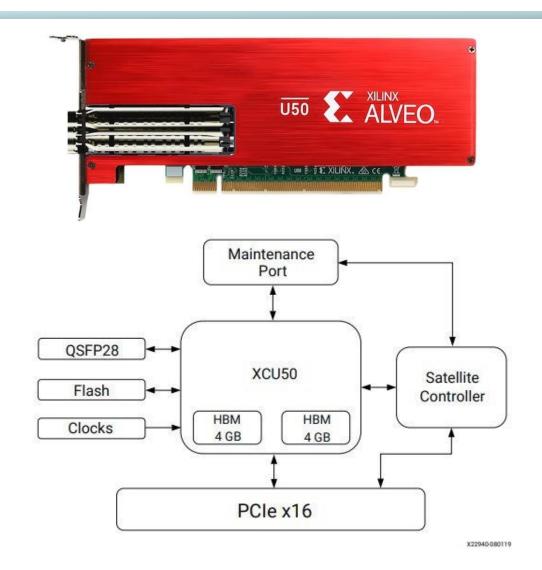




Altera Stratix/Arria 10

- > 1 million LUTs
- Floating-point cores in fabric
- HyperFlex Interconnect
 - Embeds flip flops into reconfigurable interconnect
 - Enables significantly faster clock speeds
- https://www.altera.com/products/fpga/stratix-series/stratix-10/overview.html

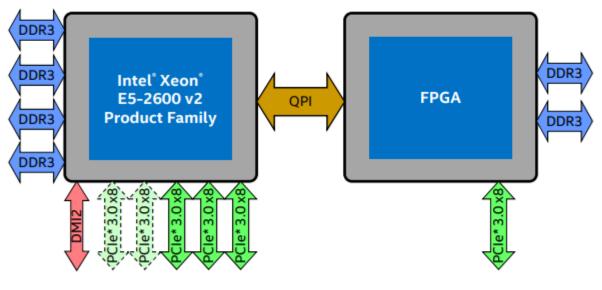
Xilinx Alveo U50 Data Center Accelerator Card



FEATURES	ALVEO U50
Architecture	UltraScale+
Form Factor	Half-Height, Half length single slot Low-Profile
Look Up Tables	872,000
HBM2 Memory	8GB
HBM2 Bandwidth	316GB/s ¹
Network Interface	1 x QSFP28 (100GbE) ²
Clock Precision	IEEE 1588
PCI Express	PCIe Gen3 x 16, dual PCIe Gen4 x 8, CCIX
Thermal Solution	Passive
Power (TDP)	75W

INTEL XEON-FPGA HYBRID CHIP





Processor	Intel® Xeon® E5-26xx v2 Processor
FPGA Module	Altera Stratix V
QPI Speed	6.4 GT/s full width (target 8.0 GT/s at full width)
Memory to FPGA Module	2 channels of DDR3 (up to 64 GB)
Expansion connector to FPGA Module	PCIe 3.0 x8 lanes - maybe used for direct I/O e.g. Ethernet
Features	Configuration Agent, Caching Agent,, (optional) Memory Controller
Software	Accelerator Abstraction Layer (AAL) runtime, drivers, sample applications

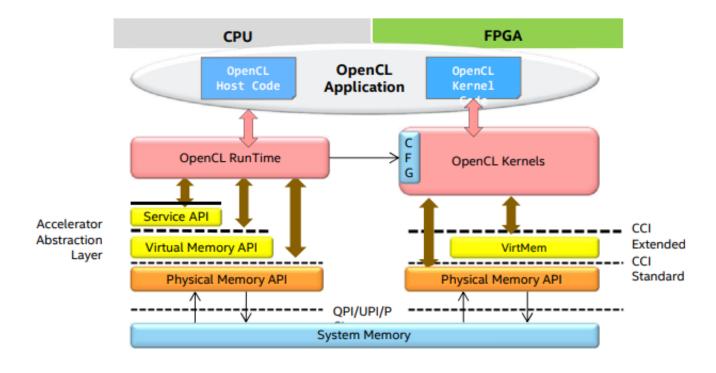


courtesy of Intel

INTEL XEON-FPGA HYBRID CHIP

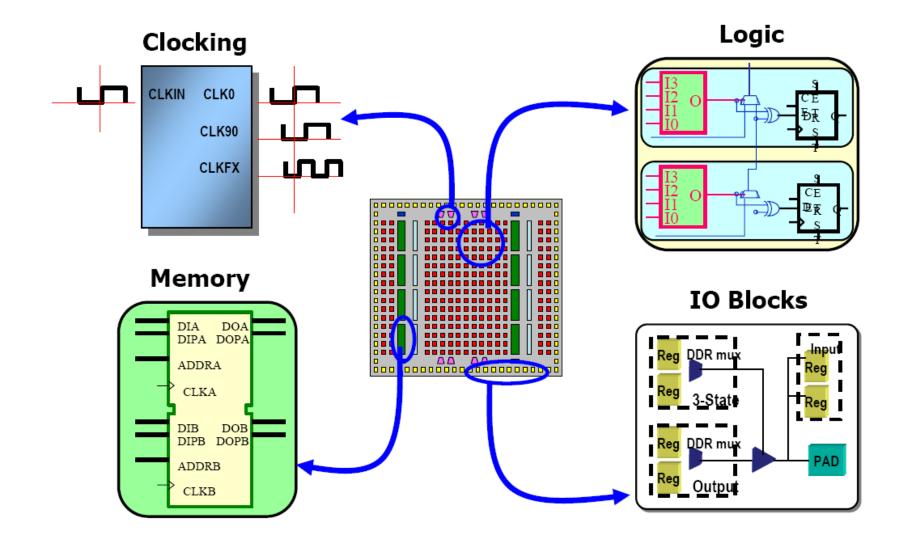


Programming Interfaces: OpenCL



courtesy of Intel

Summary



Summary

- Reconfigurable SoC is trending, FPGAs are dominant
- FPGAs
 - Better for very rapid design and redesign.
 - Good for prototyping but also end design.
 - Better for small numbers of products
 - More expensive, and less in performance
 - more efficient than GPU
 - Spatial and temporal parallelism
 - Programmable SoCs can be used as accelerators, especially in data centers

Where are we Heading?

Advanced Topics: Advanced Packaging for SoCs

Action Items

- Final projects Start ASAP
- Reading Assignment
 - Ch. 6.5, Slides

Acknowledgement

Slides in this topic are inspired in part by material developed and copyright by:

- Synopsys Courseware
- Prof. Greg Stitt (U of Florida)